

Fixed before any outcome was computed: lineage rules, caps and seed, pair model, simulation gate (three dated amendments)

1 Source

Open LLM Leaderboard, version 1
6,896 models; MMLU, 14,042 items
per-item details with option scores

2 Lineage

declared base_model chain → root
eligible: 721 primary (card links),
1,295 expanded (+ config paths)

3 Freeze

≤ 40 models per root, seed 0
613 primary + 967 expanded
(1,018 unique; SHA-256 recorded)

4 Validate

argmax of option scores must match
stored correctness on every item
977 accepted, 41 excluded (logged)

5 Pairs

$a_{ij} = P(\text{same wrong} \mid \text{both wrong})$
173,755 primary, 430,128 expanded
≥ 1,166 jointly wrong items per pair

6 Estimate

shared root + release-gap bins
+ accuracy sum and difference
12 pre-registered analyses

7 Audit (partly post hoc)

root-level inference, outcomes,
item-difficulty nulls, gate audit,
size & architecture controls