

Fixed before any pair outcome was computed: lineage rules, caps and seed, pair model, simulation gate
(three dated amendments; the third after per-model accuracies were seen)

1 Source

Open LLM Leaderboard v1
6,896 models
MMLU, 14,042 items

2 Lineage

declared base_model chain → root
721 primary (card links)
1,295 expanded (+ config paths)

3 Freeze

≤ 40 per root, seed 0, SHA-256
613 primary, 967 expanded
(overlapping: 1,018 unique)

4 Validate

argmax must match stored
correctness on every item
977 accepted (union of both),
41 excluded and logged

5 Pairs

$a_{ij} = P(\text{same wrong} \mid \text{both wrong})$
173,755 primary pairs
430,128 expanded pairs

6 Estimate

shared root + gap bins
+ accuracy terms
12 pre-specified analyses

7 Audit (partly post hoc)

root-level inference, outcomes,
item-difficulty nulls,
gate audit, size controls